

Towards Accurate Binary Convolutional Neural Network

2018-11-05

Overview

- NIPS 2017
- 做了weight和activation两者的二值化
- 1. approximating full-precision weights with the linear combination of multiple binary weight bases;
- 2. employing multiple binary activations to alleviate information loss.

ABC-Net (A ccurate-B inary-C onvolutional)

- Weight approximation

$$\min_{\alpha, \mathbf{B}} J(\alpha, \mathbf{B}) = \|\mathbf{w} - \mathbf{B}\alpha\|^2, \quad \text{s.t. } B_{ij} \in \{-1, +1\}, \quad (1)$$

where $\mathbf{B} = [\text{vec}(\mathbf{B}_1), \text{vec}(\mathbf{B}_2), \dots]$, where $\bar{\mathbf{W}} = \mathbf{W} - \text{mean}(\mathbf{W}) = [\alpha_1, \alpha_2, \dots, \alpha_M]^T$.

$$\mathbf{B}_i = F_{u_i}(\mathbf{W}) := \text{sign}(\bar{\mathbf{W}} + u_i \text{std}(\mathbf{W})), i = 1, 2, \dots, M, \quad (2)$$

$$\min_{\alpha} J(\alpha) = \|\mathbf{w} - \mathbf{B}\alpha\|^2, \quad (3)$$

- 做weight的近似组合，使用alpha逼近

ABC-Net (A ccurate-B inary-C onvolutional)

Forward: $\mathbf{B}_1, \mathbf{B}_2, \dots, \mathbf{B}_M = F_{u_1}(\mathbf{W}), F_{u_2}(\mathbf{W}), \dots, F_{u_M}(\mathbf{W})$, (4)

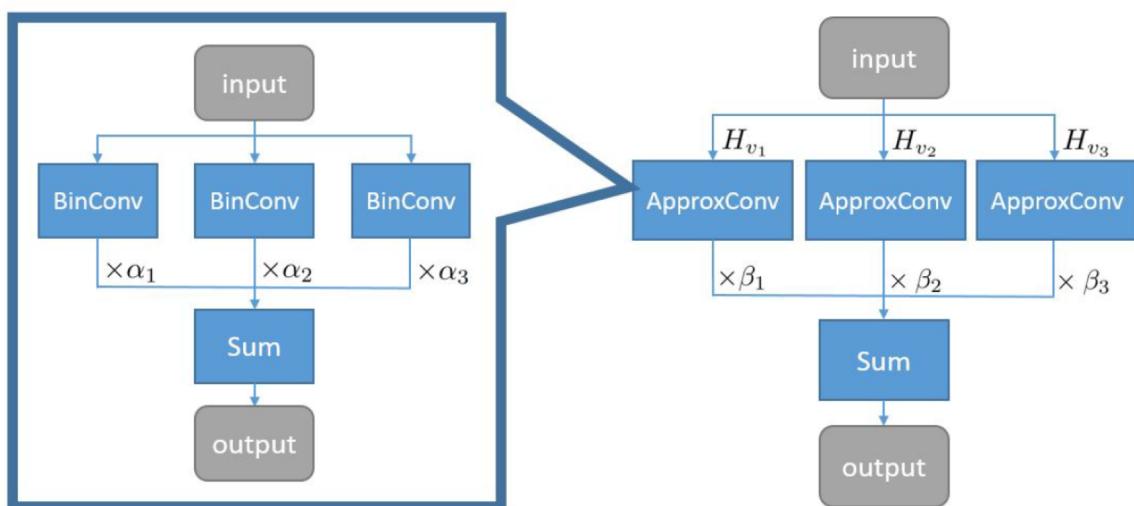
Solve (3) for α ,

$$\mathbf{O} = \sum_{m=1}^M \alpha_m \text{Conv}(\mathbf{B}_m, \mathbf{A}). \quad (6)$$

Weight部分

Backward: $\frac{\partial c}{\partial \mathbf{W}} = \frac{\partial c}{\partial \mathbf{O}} \left(\sum_{m=1}^M \alpha_m \frac{\partial \mathbf{O}}{\partial \mathbf{B}_m} \frac{\partial \mathbf{B}_m}{\partial \mathbf{W}} \right) \stackrel{\text{STE}}{=} \frac{\partial c}{\partial \mathbf{O}} \left(\sum_{m=1}^M \alpha_m \frac{\partial \mathbf{O}}{\partial \mathbf{B}_m} \right) = \sum_{m=1}^M \alpha_m \frac{\partial c}{\partial \mathbf{B}_m}. \quad (7)$

For/Back-ward



Experiment

Table 2: Prediction accuracy (Top-1/Top-5) for ImageNet with different choices of M and N in a ABC-Net (approximate weights as a whole). “res18”, “res34” and “res50” are short for Resnet-18, Resnet-34 and Resnet-50 network topology respectively. M and N refer to the number of weight bases and activations respectively.

Network	M -weight base	N -activation base	Top-1	Top-5	Top-1 gap	Top-5 gap
res18	1	1	42.7%	67.6%	26.6%	21.6%
res18	3	1	49.1%	73.8%	20.2%	15.4%
res18	3	3	61.0%	83.2%	8.3%	6.0%
res18	3	5	63.1%	84.8%	6.2%	4.4%
res18	5	1	54.1%	78.1%	15.2%	11.1%
res18	5	3	62.5%	84.2%	6.8%	5.0%
res18	5	5	65.0%	85.9%	4.3%	3.3%
res18	Full Precision		69.3%	89.2%	-	-
res34	1	1	52.4%	76.5%	20.9%	14.8%
res34	3	3	66.7%	87.4%	6.6%	3.9%
res34	5	5	68.4%	88.2%	4.9%	3.1%
res34	Full Precision		73.3%	91.3%	-	-
res50	5	5	70.1%	89.7%	6.0%	3.1%
res50	Full Precision		76.1%	92.8%	-	-

Experiment

Table 1: Top-1 (left) and Top-5 (right) accuracy of ABC-Net on ImageNet, using full-precision activation and different choices of the number of binary weight bases M .

	BWN	$M = 1$	$M = 2$	$M = 3$	$M = 5$	FP
Top-1	60.8%	62.8%	63.7%	66.2%	68.3%	69.3%
Top-5	83.0%	84.4%	85.2%	86.7%	87.9%	89.2%

Table 3: Classification test accuracy of CNNs trained on ImageNet with Resnet-18 network topology. ‘W’ and ‘A’ refer to the weight and activation bitwidth respectively.

Model	W	A	Top-1	Top-5
Full-Precision Resnet-18 [full-precision weights and activation]	32	32	69.3%	89.2%
BWN [full-precision activation] Rastegari et al. [2016]	1	32	60.8%	83.0%
DoReFa-Net [1-bit weight and 4-bit activation] Zhou et al. [2016]	1	4	59.2%	81.5%
XNOR-Net [binary weight and activation] Rastegari et al. [2016]	1	1	51.2%	73.2%
BNN [binary weight and activation] Courbariaux et al. [2016]	1	1	42.2%	67.1%
ABC-Net [5 binary weight bases, 5 binary activations]	1	1	65.0%	85.9%
ABC-Net [5 binary weight bases, full-precision activations]	1	32	68.3%	87.9%

Future Work

这一部分比较重要，硬件友好的运算优化

5.3 Further computation reduction in run-time

On specialized hardware, the following operations in our scheme can be integrated with other operations in run-time and further reduce the computation requirement.

(1) Shift operations. The existence of shift parameters seem to require extra additions/subtractions (see (2) and (8)). However, the binarization operation with a shift parameter can be implemented as a comparator where the shift parameter is the number for comparison, e.g., $H_v(\mathbf{R}) = \begin{cases} 1, & \mathbf{R} \geq 0.5 - v; \\ -1, & \mathbf{R} < 0.5 - v. \end{cases}$ ($0.5 - v$ is a constant), so no extra additions/subtractions are involved.

(2) Batch normalization. In run-time, a batch normalization is simply an affine function, say, $\text{BN}(\mathbf{R}) = a\mathbf{R} + b$, whose scale and shift parameters a, b are fixed and can be integrated with v_n 's. More specifically, a batch normalization can be integrated into a binarization operation as follow: $H_v(\text{BN}(\mathbf{R})) = \begin{cases} 1, & a\mathbf{R} + b \geq 0.5 - v; \\ -1, & a\mathbf{R} + b < 0.5 - v. \end{cases} = \begin{cases} 1, & \mathbf{R} \geq (0.5 - v - b)/a; \\ -1, & \mathbf{R} < (0.5 - v - b)/a. \end{cases}$ Therefore, there will be no extra cost for the batch normalization.